

Reproduction of vllm-project/vllm: offline throughput of Qwen2.5-1.5B-Instruct

K-Veritas Team

Submitted by K-Veritas Team

kv:2610.00030v1 · Published 2026-10-08

Independent reproduction · Artifact evaluation · Systems and performance

Abstract

Ran vLLM's offline throughput benchmark with the released vLLM 0.31.0 package (PyTorch 2.13 for CUDA 13.0, FlashInfer kernels compiled for sm_120 at startup): Qwen2.5-1.5B-Instruct on 2,000 random prompts of 512 input and 256 output tokens, default engine settings. The bundle holds the v0.31.0 tag's benchmark code and project files; the rest of the source tree was set aside because the installed package, recorded in the signed environment, is what ran. One NVIDIA RTX 5060 Ti.

About this record

Title, authors and abstract were provided by the authors and reviewed for completeness, not for scientific soundness. Every number in this record is copied from the signed data of the sealed reports [2], which remain authoritative.

This cover page is not signed. Check it against the SHA-256 listed at <https://kveritas.org/records/2610.00030v1>, and verify each report with its code bundle: `kveritas verify <report.pdf> --bundle <bundle.zip>`.

1 Sealed reports

The experiments were sealed with K-Veritas [1], which binds each result to the code, hardware and time that produced it.

Report	Label	Sealed	Runs	Invocations	Signer
r1 [2]	Qwen2.5-1.5B throughput, 20...	2026-10-08	1	1	Signed by the K-Veritas server

2 Results

r1: Qwen2.5-1.5B throughput, 2000 prompts

Run	Metric	Final value	Step
1	elapsed_seconds	193.688	
1	num_requests	2000	

Run	Inline claim	Value	Phase
1	requests_per_second	10.3259	
1	tokens_per_second	11895.4	

3 Setup

Report	Item	Value
r1	Operating system	linux / amd64
r1	CPU	Intel(R) Xeon(R) W-2145 CPU @ 3.70GHz, 16 cores
r1	GPU	NVIDIA GeForce RTX 5060 Ti, 16311 MiB
r1	Memory	31.0 GB
r1	Total run time	247.0 s

4 Files

Each report and code bundle must hash to these values. A bundle belongs to the report whose signed data names its hash.

File	SHA-256	Size
r1.pdf	83441ef3015e5de579d5b393051b13058ed45c12eabc010854860584657e2cb8	2.4 MB
r1.kvbundle.zip	dd5a8c661b389af8c024db7f4467231a74a33139ded7905516ebd032890aba8d	647 KB

5 Links and licence

Paper: <https://arxiv.org/abs/2309.06180>

Reproduction of: <https://github.com/vllm-project/vllm/tree/db9527a46873454610df6dbedf79a36d6bf1a7f6>

Record: CC-BY 4.0. Code: Apache-2.0, as declared by the authors. Tags: llm-inference, vllm, throughput, paged-attention.

Listed, not endorsed. Review checks the artifact, never the science. Reports are never edited; changing them makes a new version. Other fields are corrected in place, and every correction is logged in metadata.yaml.

References

[1] M. K. Keita and C. Homan. Computer Science Conferences Should Require Nonrepudiable Experimental Results. NeurIPS 2026 Position Paper Track. arXiv:2605.08586.

[2] Sealed experiment report r1 (Qwen2.5-1.5B throughput, 2000 prompts), kv:2610.00030v1/r1. Data hash bc31cd8c61d010a64596faef5b0e1343ad41bb2c43fb04c836c027879dd8b50f.

<https://kveritas.org/records/2610.00030v1/r1.pdf>